

Modularity for Decidability of Deductive Verification with Applications to Distributed Systems

Anonymous Author(s)

Abstract

Proof automation can substantially increase productivity in the formal verification of complex systems. However, the unpredictability of automated provers in handling quantified formulas presents a major hurdle to usability of these tools. Here, we propose to solve this problem not by improving the provers, but by using a modular proof methodology that allows us to produce *decidable* verification conditions. Decidability greatly improves the predictability of proof automation, resulting in a more practical verification approach. We apply this methodology to develop verified implementations of distributed protocols with competitive performance, showing its effectiveness.

1 Introduction

Verifying complex software systems is a longstanding research goal. Recently there have been some success stories in verifying compilers [26], operating systems [19], and distributed systems [15, 40]. These broadly use two techniques: interactive theorem proving (e.g., Coq [4], Isabelle/HOL [33]) and deductive verification based on automated theorem provers (e.g., Dafny [25] which uses Z3 [10]). However, both techniques are difficult to apply and require a large proof engineering effort. On the one hand, interactive theorem provers allow a user to write proofs in highly expressive formalisms (e.g., higher-order logic or dependent type theory). While this allows great flexibility, it generally requires the user to manually write long and detailed proofs.

On the other hand, deductive verification techniques use automated theorem provers to reduce the size of the manually written proofs. In this approach, user-provided annotations (e.g. invariants, pre- and post-conditions) are used to reduce the proof to lemmas called *verification conditions* that can be discharged by the automated prover. In case these lemmas fail, the prover can sometimes produce counterexamples that explain the failure and allow the programmer to correct the annotations.

Unfortunately, the behavior of provers can be quite unpredictable, especially when applied to formulas with quantifiers, which are common in practice, e.g. in distributed systems. Since the problem presented to the prover is in general undecidable, it is no surprise that the prover sometimes diverges or produces inconclusive results on small instances,

or suffers from the “butterfly effect”, when a seemingly irrelevant change in the input causes the prover to fail. As observed in the IronFleet project, SMT solvers can diverge even on tiny examples [15]. When this happens, the user has little information with which to decide how to proceed. This was identified in IronFleet as the main hurdle in the verification task.

One approach to address the unpredictability of automated solvers is to restrict verification conditions to a decidable fragment of logic [3, 12, 16, 30]. Our previous work on the Ivy verification system [31, 35, 36] has used the effectively propositional (EPR) fragment of first-order logic to verify distributed protocols designs (e.g. cache coherence and consensus). However, the restrictions imposed for decidability are a major limitation, and in particular, limit the applicability of previous work to protocol designs and not their executable implementations. I found this hard to parse. "(ap. of p.w.) to p.d." or "ap. of (p.w. to p.d.)?"

Decidable Decomposition In this paper, we show how to use well-understood modular reasoning techniques to increase the applicability of decidable reasoning and support verifying implementations as well as designs. The key idea is to structure the correctness proof in a modular way, such that each component can be proved using a decidable fragment of first-order logic, possibly with a background theory. Importantly, each component’s verification condition can use a *different* decidable fragment. This allows, for example, one component to use arithmetic, while another uses stratified quantifiers and uninterpreted relations. It also allows each component to use its own quantifier stratification, even when the combination would not be stratified. We will refer to this approach as *decidable decomposition*. Crucially, decidable decomposition can be applied even when the global verification condition does not lie in any single decidable fragment (for example, the combination of arithmetic, uninterpreted relations, and quantifiers is undecidable).

Because our prover is a decision procedure for the logical fragments we use, we can guarantee that in principle it will always terminate with a proof or a counter-model. In practice, decidability means that the behavior of the prover is much more predictable.

Verifying Distributed Systems As a demonstration of decidable decomposition, we verify distributed protocols and their implementations. Distributed protocols play an essential role in today’s computing landscape, but our previous work on the Ivy verification system could not verify their implementations. In particular, reasoning about distributed

111 protocols naturally leads to quantifiers and uninterpreted
 112 relations, while their implementations use both arithmetic
 113 and concrete representations (e.g., arrays). This combination
 114 escapes the decidable fragments used in previous work.
 115

We observe in our evaluation that the human effort needed to achieve decidable decomposition is modest. We identify a few patterns that suffice to decompose interesting systems. For example, in our implementation of Multi-Paxos, we decompose the proof into an abstract protocol and an implementation, where each component's verification condition falls in a (different) decidable fragment. At the end of the process, we compile the verified system to executable code (which uses a small set of trusted libraries, e.g., to implement a built-in array type). Our preliminary experience indicates that this can be used to generate reasonably efficient verified distributed systems.

Contributions The contributions of this paper are:

1. A new methodology, *decidable decomposition*, that uses existing modularity principles to decompose the verification into lemmas proved in (different) decidable logics.
2. A realization of this methodology in a deductive verification tool, Ivy, that supports compilation to C++ and discharges verification conditions using an SMT solver. The fact that all verification conditions are decidable makes the SMT solver's performance more predictable, improving the system's usability and reducing verification effort.
3. An application of the methodology to distributed systems, resulting in verified implementations of two popular distributed algorithms, Raft and Multi-Paxos, obtaining reasonable run-time performance. We show that proofs of these systems naturally decompose into decidable sub-problems. Our experience is that verifying systems in decidable logics is significantly easier than previous approaches.

2 Overview

In this section, we motivate and demonstrate our key ideas on a simple example.

2.1 Example: Toy Leader Election

Figure 1 shows pseudocode for a node that participates in a toy leader election protocol, in which a finite set of nodes decide on a leader. The set of nodes is a parameter of the system, which is determined at run time and remains fixed throughout each run of the protocol. Each node may propose itself as a candidate by sending a message to all nodes. Nodes vote, by sending a response message, for the first candidate from which they receive a message. A leader is elected when it receives a majority of the votes. This protocol will get stuck in many cases without electing a leader. However, it suffices to demonstrate our verification methodology, since

This code is still quite opaque, in terms of semantics of message send and receive, "self", etc.

```

1 // spec: at most one node           11 upon recv(msg) do {
2 //     sends leader_msg            12   if msg.type = request_vote_msg
3                                13     ^-alreadyvoted {
4 alreadyvoted := false          14       alreadyvoted := true;
5 voters := ∅                      15       send vote_msg(self, msg.src)
6 upon client_request() do {      16     } else if msg.type = vote_msg {
7   if ¬alreadyvoted {             17       voters := voters ∪ { msg.src }
8     send request_vote_msg(self) 18       if |voters| > N/2 {send leader_msg(self)}
9   }                           19   }
10 }                            20 }
```

Figure 1. Toy Leader Election pseudocode.

it is *safe*, i.e., at most one leader is elected. Furthermore, a variant of this protocol is an essential ingredient of both Raft and Multi-Paxos, which is used in many production systems, as well as in our evaluation.

The goal of the verification is to show that at most one leader is elected. Despite the simplicity of the property and the code, existing verification techniques cannot automatically prove that the code is correct when executed by an unbounded number of nodes. Even when the code is annotated with invariants, the corresponding verification conditions are expressed in undecidable logics. As a result, checking them with existing theorem provers such as Z3 [10] often results in divergence, and behaves unpredictably in general. Indeed, previous verification efforts in the systems community identified this as a major hurdle for verification [15]. The complexity arises due to the combination of arithmetic, set cardinalities (e.g., number of nodes that voted for a candidate), and quantifiers in the invariant that quantify over unbounded domains (e.g., expressing the fact that for every two nodes at most one is a leader), especially *non-stratified* quantifier alternations (see Section 3.2), which give rise to potentially infinitely many instantiations.

2.2 Approach

In this paper we present a verification methodology based on decidable reasoning. We use it to develop and verify implementations of distributed systems, whose performance is similar to other verified implementations (e.g. [41]). Rather than starting with an existing implementation (e.g., in C), we define a simple imperative language which permits effective (decidable) reasoning on one hand, and straightforward compilation to efficient C++ code on the other hand.

Our method leverages modularity in the assume-guarantee style for decidable reasoning, by structuring the correctness proof such that different parts of the proof reason about different aspects of the system, and at different abstraction levels, enabling each of them to be carried in (possibly different) decidable logics. The decomposition has two benefits. First, it allows to reduce quantifier alternations, and eliminate bad quantification cycles. Second, it allows to check the verification conditions of each module using a different background theory (or none).

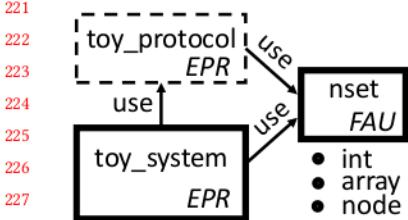


Figure 2. Modules and built-in types used for verifying the toy example. Dashed box denotes a ghost module. Each module is annotated with the decidable fragment in which it is verified.

2.3 Modular Formulation

We illustrate our methodology on the Toy Leader Election example, to verify that at most one leader is elected using *decidable* reasoning. Our formulation of the system consists of three *modules*: *toy_protocol*, *toy_system*, and *nset*. The interplay between the modules and the decidable fragments in which they are verified are depicted in Figure 2 (the fragments are defined in Section 3); their code is listed in Figures 3 to 5. The code is written in **Modular Decidable Language** (MDL) — an illustrative programming language enabling modular decidable verification.

Each module should be viewed as a proof unit, which consists of (1) declarations and definitions of types and state components, which may be interpreted using **interpret** declarations, (2) declarations of other modules and their invariants that are used by the module, specified by the **uses** clause, (3) a module invariant *Q*, given by all **invariant** declarations in the module, and (4) procedures with pre-post specifications, specified by **requires** and **ensures** declarations (an unspecified condition is *true* by default), (5) declaration of the module's initial state, either with **init** declarations or with an **init()** procedure. Intuitively, a module is correct if all its procedures satisfy their pre/post specification, and also maintain the module invariant, assuming that the used modules are themselves correct. The key property of the modular formulation is that verification conditions generated for each module are decidable to check (see Section 2.4), which allows predictable automation. better: "fall in dec'l theories"

We elaborate on the modules of the Toy Leader Election example in Sections 2.3.1 to 2.3.3. Roughly speaking:

(i) The *nset* module defines (and verifies) a data type which encapsulates sets of nodes under a first-order interface, and hides the low-level implementation. This allows other modules to treat sets of nodes as an opaque type, relying only on the first-order interface, so their verification can be carried out in uninterpreted first-order logic. Verifying that the *nset* module satisfies its interface is carried out in a suitable theory. The first-order interface includes a predicate that tests if a set of nodes forms a majority, along with the property that any two majorities intersect, which is crucial for the proof of the protocol.

(ii) The *toy_protocol* module defines (and verifies) an abstract model of the protocol, eliding some of the implementation details. This is in line with the common practice of

```

1 ghost module toy_protocol uses nset, nset.majorities_intersect {           276
2   relation voted : node, node                                         277
3   relation isleader : node                                         278
4   variable quorum : nset.t                                         279
5   init ∀n1, n2. ¬voted(n1, n2)                         280
6   init ∀n. ¬isleader(n)                                         281
7   invariant one_leader = ∀n1, n2. isleader(n1) ∧ isleader(n2) → n1 = n2 282
8   invariant ∀n, n1, n2. voted(n, n1) ∧ voted(n, n2) → n1 = n2 283
9   invariant ∀n : node. isleader(n) → (nset.majority(quorum) ∧          284
10                                     ∀n' : node. nset.member(n', quorum) → voted(n', n)) 285
11  procedure vote(v : node, n : node) {                                286
12    requires ∀n' .¬voted(v, n')                                         287
13    voted(v, n) := true                                              288
14  }                                                               289
15  procedure become_leader(n : node, s : nset.t) {                      290
16    requires nset.majority(s) ∧ ∀n' : node. nset.member(n', s) → voted(n', n) 291
17    isleader(n) := true                                              292
18    quorum := s                                              293
19  }                                                               294
20  }                                                               295

```

Figure 3. Protocol module for Toy Leader Election.

developing a system in an evolutionary process: starting with a design, and then gradually developing an efficient implementation. The *toy_protocol* captures the protocol design and its correctness in a ghost object, which is verified separately and serves as a lemma for proving the implementation. This provides a natural way to decompose the verification problem, and as we shall see, to avoid quantifier alternation cycles.

(iii) The *toy_system* module specifies (and verifies) the system implementation, using both the data type defined by *nset* (relying on its first-order specification) and the abstract protocol defined by *toy_protocol* (as ghost code) to obtain a verified executable implementation.

2.3.1 Abstract Protocol Module

Figure 3 lists the module that formalizes the abstract leader election protocol. This is a ghost module, which is only used for the sake of the proof. The module contains two mutable relations that define its state: *voted(n₁, n₂)* captures the fact that *n₁* voted for *n₂*, and *isleader(n)* means node *n* is an elected leader. The initial state of the module specifies that both relations are empty. The state also includes a variable *quorum*, that remembers the last voting majority observed. The abstract protocol provides a global invariant (denoted by **invariant**), which states that there is at most one leader. This is similar to a class invariant / object invariant in modular reasoning. Next, the module contains a proof that all of its reachable states satisfy the global invariant, by an inductive invariant. When proving this module, the majority intersection invariant of the *nset* module is used, as indicated by the **uses** clause in line 1). ↗ unbalanced paren

The module also provides two procedures that define the abstract protocol steps. Each procedure specifies a pre condition. The *vote(n₁, n₂)* procedure models a vote by *n₁* for

```

331 1 system module toy_system uses nset, toy_protocol, toy_protocol.one_leader {
332 2   message request_vote_msg : node
333 3   message vote_msg : node, node
334 4   message leader_msg : node
335 5   // spec: at most one node sends leader_msg :
336 6   invariant safe = ∀n1, n2. leader_msg(n1) ∧ leader_msg(n2) → n1 = n2
337 7
338 8   relation alreadyvoted : node
339 9   function voters : node → nset.t
340 10  procedure init(self : node) {
341 11    alreadyvoted(self) := false
342 12    voters(self) := nset.emptyset()
343 13  }
344 14  procedure request_vote(self : node) {
345 15    send request_vote_msg(self)
346 16  }
347 17  procedure cast_vote(self : node, n : node) handles request_vote_msg(n) {
348 18    if ¬alreadyvoted(self) {
349 19      alreadyvoted(self) := true
350 20      send vote_msg(self, n)
351 21      toy_protocol.vote(self, n)
352 22    }
353 23  }
354 24  procedure receive_vote(self : node, n : node) handles vote_msg(n, self) {
355 25    voters(self) := nset.add(voters(self), n)
356 26    if nset.majority(voters(self)) {
357 27      send leader_msg(self)
358 28      toy_protocol.become_leader(self, voters(self))
359 29    }
360 30  }
361 31 // inductive invariant for the proof:
362 32 invariant ∀n1, n2. toy_protocol.voted(n1, n2) ↔ vote_msg(n1, n2)
363 33 invariant ∀n1, n2. nset.member(n1, voters(n2)) → toy_protocol.voted(n1, n2)
364 34 invariant ∀n. leader_msg(n) ↔ toy_protocol.isleader(n)
365 35 invariant ∀n1, n2. ¬alreadyvoted(n1) → ¬toy_protocol.voted(n1, n2)
366 36 open toy_protocol
367 37 }
```

Figure 4. System module for Toy Leader Election.

*n*₂, and its precondition is that node *n*₁ has not yet voted. The *become_leader*(*n*, *s*) procedure models the election of *n* as a leader, and its precondition requires that all nodes in *s* voted for *n*, and that *s* is a majority. Note that this module is abstract in the sense that it abstracts network communication and uses a global view of the system. In particular, the *become_leader* does not specify how a node learns that it received a majority of votes.

2.3.2 Concrete Implementation Module

Figure 4 lists the concrete system implementation. We consider systems in which a finite (but unbounded) set of nodes run the same code, and exchange messages. For simplicity, we assume all messages are broadcast to all nodes. Our network model also allows message dropping, duplication and reordering. To define the system implementation, we first define the message types. Lines 2 to 4 define three message types: *request_vote_msg* that a node uses to propose itself as a leader, *vote_msg* that a node sends to vote for a candidate, and *leader_msg* that a node uses to announce it is elected as the leader. The first field of each message is its source node. The invariant on line 6 expresses the specification that

that only one node ever issues a *leader_msg* message. This is the unique guarantee provided by the implementation, and as such it is the only line of trusted code in the example. These declarations are followed by the code to be run on each node. This code (lines 8 to 30) defines the local state of each node, as well as procedures that can be executed in response to client requests, or procedures that are message handlers (specified by the **handles** declaration), and executed upon receiving a message from the network. The state components are defined as functions (or relations) whose first argument is a node, so that *f*(*n*) denotes the local state of node *n*. Similarly, all procedures receive as their first argument the *self* identifier of the node that runs them. Since the system module is going to be compiled into an executable code that runs on each node, we syntactically enforce that each node may only access and modify its own local state.

We note that the *toy_system* module makes use of the *nset* module to maintain sets of voters, and uses a majority test (line 26). It also makes calls into the ghost module *toy_protocol* (lines 21 and 28). This allows to establish an invariant that relates the state of the concrete module and the abstract module (lines 32 to 35), and use the proven invariant of the abstract module as a lemma for proving the concrete implementation. This is indicated by the **uses** clause, which declares use of the *toy_protocol.one_leader* invariant (line 1).

2.3.3 Node Set Module

Figure 5 lists the code for the *nset* module. This module defines a data structure for storing sets of nodes, with operations for adding to a set and testing whether a set is a majority. To do so, it defines a type *t*, whose internal interpretation is MDL's built-in array, as declared in line 7. Sets are created using the *emptyset* and *add* procedures, which provide a naive implementation of a set, stored as an array of its elements. The module defines the *member* relation, and provides it with an interpretation via an *interpret* declaration in line 8. As we shall see, this definition creates an executable membership test that is translated to a loop that scans the array.

Most importantly for verifying the leader election modules, the *nset* module provides the interpreted majority predicate, with the key property that any two majorities intersect. This is stated by the *majorities_intersect* invariant (line 11). Intuitively, a set is a majority if its cardinality is more than half the total number of nodes. In Figure 5, the majority predicate is interpreted using the *card* function to compute the cardinality of a set of nodes, and the built-in value *node.all*, which is an array with the special semantics that it contains all nodes. The *card* function computes the cardinality of a set of nodes, and it is constructed in the *init()* procedure of the *nset* module in a way that establishes the majority intersection property. The proof is by induction, manifested in the loop invariant in line 25.

```

441 1 module nset {
442   2 type t
443   3 relation member : node, t
444   4 relation majority : t
445   5 function card : t → int
446   6
447   7 interpret t as array<int,node>
448   8 interpret member(n, s) as ∃i. 0 ≤ i < array.len(s) ∧ array.value(s, i, n)
449   9 interpret majority(s) as card(s) + card(s) > card(node.all)
450  10
451  11 invariant majorities_intersect = ∀s1, s2. majority(s1) ∧ majority(s2) →
452    12   ∃n. member(n, s1) ∧ member(n, s2)
453
454  13
455  14 procedure emptyset() returns s:t {
456    15   ensures ∀n. ¬member(n, s)
457    16   s := array.empty()
458  17 }
459  18 procedure add(s1 : t, n : node) returns s2 : t {
460    19   ensures ∀n'. member(n', s2) ↔ (member(n', s1) ∨ n' = n)
461    20   if member(n, s1) then { s2 := s1 } else { s2 := array.append(s1, n) }
462  21 }
463  22 procedure init() {
464    23   card := λx. 0;
465    24   for 0 ≤ i < array.len(node.all) {
466     25     invariant ∀s1, s2. (∀n. ¬(member(n, s1) ∧ member(n, s2))) →
467       26       card(s1) + card(s2) ≤ card(node.all)
468     27     card := λx. ((card(x) + 1) if member(array.get(node.all, i), x) else card(x))
469   28   }
470  29 }
471  30 } This part is quite subtle, apparently reasoning about
472    closures despite f.o. logic! Is there more worth saying?

```

Figure 5. nset module for node sets, proving the majority intersection property.

We note that the loop in init() constructs card via a nesting of N function closures (where N is the number of nodes). This definition of card allows an easy proof of the majority intersection property. A more efficient implementation uses array.len (underlying array length) instead of card to determine if a set is a majority. This implementation is also provable in our system, but requires additional inductive invariants to prove that card and array.len coincide, and we do not present it here in the interest of simplicity.

2.4 Modular Verification in Decidable Fragments

We now explain how verification conditions are generated for each module, and how they are checked under (possibly different) theories. We use two decidable fragments: the effectively propositional (EPR) fragment [37] which allows stratified quantifier alternations, and the finite almost uninterpreted (FAU) fragment [14] which includes linear integer arithmetic in a restricted way.

Verification Condition Generation Recall that each module declares the modules and invariants it uses in the **uses** clause. Based on this, verification conditions are derived automatically. Each module must provide the following *guarantees*: 1) The module invariant, Q , holds at any initial state. 2) Every procedure in the module establishes its postcondition and preserves the module invariant Q , and

3) At each call site, the precondition of the called procedure is established. Each module may rely upon the following *assumptions*: 1) Every called procedure establishes its postcondition and 2) Every used invariant of another module holds at all times. The verification condition states that the assumptions imply the guarantees.

As an example, the verification condition generated for the become_leader procedure (Figure 3 line 15) is:

$$(Q_{\text{toy_protocol}} \wedge Q_{\text{nset.majorities_intersect}} \wedge \\ (\text{majority}(s) \wedge \forall n'. \text{member}(n', s) \rightarrow \text{voted}(n', n))) \rightarrow \\ Q_{\text{toy_protocol}} [(isleader}(x) \vee x = n) / isleader(x), s / quorum]$$

where $Q_{\text{toy_protocol}}$ is given by Figure 3 lines 7 to 9 and $Q_{\text{nset.majorities_intersect}}$ is given by Figure 5 line 11. The latter is used because it appears in the **uses** clause of the toy_protocol module. The verification condition also includes the procedure precondition taken from Figure 3 line 16, and checks that assuming the invariants and the precondition, the invariant $Q_{\text{toy_protocol}}$ is preserved by the procedure (the substitutions reflect the assignments of lines 17 and 18).

Optionally, calls to procedures may be inlined (i.e., replaced by the procedure body), and we may take the initial condition of another module as an assumption. If we use a module in this way, we say the module is *opened*. As an example, ghost module toy_protocol is opened in the verification of toy_system (fig. 4, line 36). This allows us to establish easily an invariant relating the states of the two modules. When composing modules, there are additional conditions required for soundness (e.g., that modules do not interfere), which are described in Section 4.

Theory Abstractions Recall that every module may include interpreted types (e.g., int, array) as well as definitions via **interpret** declarations. These allow the module to define its relevant *theory*. A theory is a (possibly infinite) set of first-order formulas, which may either be given explicitly (e.g., the theory of total order), or by using a built-in theory (e.g., the theory of linear integer arithmetic). The verification conditions of the module are checked with respect to the provided theory and definitions. Symbols that are given no interpretation in the module are treated as uninterpreted, which may be viewed as a form of abstraction.

For example, when checking the verification condition of the become_leader procedure given above, the majority and member relations are treated as uninterpreted relations, and not according to their definitions from Figure 5 lines 8 and 9. In contrast, when verifying the nset module, these definitions are used as part of a background theory, which also includes linear integer arithmetic.

Decidable Decomposition of Toy Leader Election For the whole picture of our example, observe that the nset module uses the int type, which introduces the theory of linear

551 integer arithmetic. Furthermore, its majorities_intersect invariant introduces quantifier alternation from nset.t to node.
 552 The function voters of the toy_system module introduces a dependency in the opposite direction, as it is a function
 553 from node to nset.t. As a result, had the nset module and its
 554 invariants been inlined within the proof of toy_system, they
 555 would have resulted in verification conditions that use the
 556 theory of arithmetic and also contain quantifier alternation
 557 cycles, breaking decidability.

558 In order to break the bad cycle, we do not directly use
 559 the majorities_intersect invariant of nset in toy_system. Instead,
 560 our proof exploits the toy_protocol module and its
 561 one_leader invariant (which does not introduce dependency
 562 from node to nset.t). Namely, the invariant one_leader is
 563 assumed when verifying toy_system, and is verified sep-
 564 arately as part of the toy_protocol module (assuming the
 565 majorities_intersect invariant of nset). In this way, toy_system
 566 only contains functions from node to nset.t, while toy_protocol
 567 only contains dependencies from nset.t to node, avoiding
 568 quantifier alternation cycles.

569 In terms of theories, the nset module is verified when int is
 570 interpreted using the theory of linear integer arithmetic, and
 571 the nodeset.t type is interpreted as array<int, node> from
 572 our built-in array theory (as explained in Section 4, we en-
 573 code arrays in FAU). Moreover, the member and majority
 574 relations are interpreted by their definitions (Figure 5 lines 8
 575 and 9). The resulting verification conditions for this module
 576 are in FAU. When verifying the toy_protocol and toy_system
 577 modules, all sorts and relations (including member and majority)
 578 are uninterpreted. The resulting verification conditions are
 579 in EPR, as the quantifiers in each module are stratified.

580 We thus see that the separation between the three modules
 581 allows us to obtain decidable verification conditions.

582 2.5 Compiling to C++ and Runtime System

583 In order to obtain a verified implementation, Ivy generates
 584 C++ code. During this phase, ghost code (the ghost module
 585 toy_protocol in our example) is sliced out, and every call to a
 586 procedure of a ghost module is treated as skip. The remaining
 587 code is translated to C++, as detailed next.

588 **Translation of Primitive Language Constructs** Every
 589 procedure is translated to a C++ function in a straightforward
 590 syntax-directed manner. Control flow constructs are trans-
 591 lated into the corresponding C++ constructs. Interpreted
 592 sorts are given appropriate representations. The built-in
 593 type int is represented by machine integers, arrays are repre-
 594 sented by the STL std::vector template, and record types
 595 are represented by C++ struct.¹ Variables of function sort
 596 are represented by function closures equipped with a memo-
 597 table. Every type in a non-ghost module must have one of
 598 the above as its interpretation.

600 ¹In the current implementation, integer overflow is not addressed. We intend
 601 this to be handled by an efficient bignum package.

603 Ivy code may contain quantified formulas as control flow
 604 conditions (e.g., as the condition of an if statement). In non-
 605 ghost context, these must be of the form $\exists i : \text{int. } a \leq i \leq$
 606 $b \wedge \varphi(i)$ or $\forall i : \text{int. } a \leq i \leq b \rightarrow \varphi(i)$. Ivy translates such
 607 conditions into for loops. In the Toy Leader Election example,
 608 this mechanism is used to compile the definition of member
 609 (Figure 5 line 8) to an executable test.

610 **Network & Runtime** The generated C++ code is intended
 611 to operate in a distributed setting, where each node runs
 612 the same program, and nodes communicate via message
 613 passing (for simplicity, we assume messages are broadcast
 614 to all nodes). Accordingly, a system module must define the
 615 message types, and the local state and procedures of each
 616 node. The local state relations and functions, and the local
 617 procedures, all have an argument of the built-in type node
 618 as their first parameter, which represents the local node. The
 619 generated C++ code includes variables of the appropriate
 620 type that represent the local state of the node. It also includes
 621 the local procedures which can only access the local state of
 622 the node, and may also send messages. Some procedures are
 623 designated as message handlers.

624 The generated C++ code is linked to client code to form the
 625 complete application. The client code may call the generated
 626 procedures (such as request_vote in the example) in order
 627 to use the service provided by the verified code.

628 The generated code also includes an additional shim that
 629 takes care of sending and receiving messages, and firing
 630 timers. Namely, message sending is translated to calls to
 631 appropriate shim functions, and the shim calls message han-
 632 dlers or timeout handlers when messages are received or
 633 timers expire, respectively. The shim also initializes the val-
 634 ues of self and node.all with a node identifier, and an array
 635 of all node identifiers, respectively. This information is ob-
 636 tained at run time from a configuration file or command
 637 line arguments. The user is trusted to properly configure the
 638 system, i.e., to run processes with unique id's, and to provide
 639 each process with a correct list of all other process id's and
 640 network information (e.g., IP addresses and ports).

3 Preliminaries

641 In this section we provide a brief background on first-order
 642 logic, theories and the decidable fragments used in this work.

3.1 Formulas and theories

643 We consider many-sorted first-order logic with equality,
 644 where formulas are defined over a set S of first-order sorts,
 645 and a vocabulary Σ which consists of sorted constant sym-
 646 bols and function symbols. Constants have first-order sorts
 647 in S , while functions have sorts of the form $(\sigma_1 \times \dots \times \sigma_n) \rightarrow \tau$,
 648 where $\sigma_i, \tau \in S$. In other words, functions may not be higher-
 649 order. We assume that S contains a sort \mathbb{B} that is the sort of
 650 propositions. A function whose range is \mathbb{B} is called a relation.
 651 If s is a symbol or term and t is a sort, then $s : t$ represents the

652 653 654 655 656 657 658 659 660

constraint that s has sort t . We elide sort constraints when they can be inferred. A Σ -structure maps each sort $t \in \mathcal{S}$ to a non-empty set called the *universe* of t , and each symbol $s : t$ in Σ to a value of sort t .

For a set of formulas T , we denote by $\Sigma(T)$ the vocabulary of T , that is, the subset of Σ that occurs in T . A *theory* T is a (possibly infinite) set of formulas. We use theories to give concrete interpretations of the symbols in Σ . For example, a given sort might satisfy the theory of linear order, or the theory of arithmetic. In particular:

Definition 3.1. A theory T is *\mathcal{V} -conservative*, where $\mathcal{V} \subset \Sigma$, if every $(\Sigma \setminus \mathcal{V})$ -structure σ can be extended to a Σ -structure σ' such that $\sigma' \models T$.

Intuitively speaking, \mathcal{V} -conservativeness means that T defines the symbols in \mathcal{V} , possibly in terms of other symbols. We will also say that a \mathcal{V} -conservative theory *interprets* the symbols in \mathcal{V} . We can compose conservative theories to obtain conservative theories:

Theorem 3.2. If T is a \mathcal{V} -conservative theory and T' is a \mathcal{V}' -conservative theory, where $\Sigma(T)$ and \mathcal{V}' are disjoint, then $T \cup T'$ is a $(\mathcal{V} \cup \mathcal{V}')$ -conservative theory.

That is, definitions can be combined sequentially, provided earlier definitions do not depend on later definitions.

3.2 Decidable fragments

We consider fragments of first-order logic for which checking *satisfiability*, or *satisfiability modulo a theory* (i.e., satisfiability restricted to models of a given theory T), is decidable.

Effectively Propositional Logic (EPR) The effectively propositional (EPR) fragment of first-order logic, also known as the Bernays-Schönfinkel-Ramsey class is restricted to first-order formulas with a quantifier prefix $\exists^* \forall^*$ in prenex normal form defined over a vocabulary Σ that contains only constant and relation symbols, and where all sorts and symbols are uninterpreted. Satisfiability of EPR formulas is decidable [27]. Moreover, formulas in this fragment enjoy the *finite model property*, meaning that a satisfiable formula is guaranteed to have a finite model.

A straightforward extension of this fragment allows *stratified* function symbols and quantifier alternation, as formalized next. The *Skolem normal form* of a formula is an equisatisfiable formula in \forall^* prenex normal form that is obtained by converting all existential quantifiers to Skolem functions. The *function graph* of a formula is a graph whose vertex set is $\mathcal{S} \setminus \{\mathbb{B}\}$, having an edge (s, t) if there is a function symbol occurring in the formula's Skolem normal form with s in its domain and t in its range. Notice this means that a formula of the form $\forall x : s. \exists y : t. \varphi$ has an edge (s, t) in its function graph, since the Skolem function for y is of sort $s \rightarrow t$. A *bad cycle* in the function graph of φ is one containing a sort s such that some variable of sort s is universally quantified in the Skolem normal form of φ .

A formula is *stratified* if its function graph has no bad cycles. Notice that all EPR formulas are stratified (since all the Skolem symbols are constants) and so are all the quantifier-free formulas. A formula φ is *virtually stratified* if there is any consistent assignment of sorts to symbols in $\Sigma(\varphi)$ under which φ is stratified.

As an example, the formula $\forall x : s. \exists y : t. f(x) = y$ is stratified, since the function graph contains only the edge (s, t) . On the other hand, $\forall x : s. \exists y : t. f(y) = x$ is not stratified, because the Skolem function for y has sort $s \rightarrow t$, while f has sort $t \rightarrow s$. The formula $\forall x : s. f(g(x)) = y$ is not stratified, since g has sort $s \rightarrow t$ while f has sort $t \rightarrow s$. However, it is virtually stratified, since we can resort it as $\forall x : s. f(g(x)) = y : u$. Also, notice that $\forall x : s. f(g(x)) = y : t$ is stratified even though there is a cycle containing sort t , because this cycle does not contain a universally quantified variable.

The extended EPR fragment consists of all virtually stratified formulas. The extension maintains both the finite model property and the decidability of the satisfiability problem (this is a special case of Proposition 2 in [14]).

Finite Almost Uninterpreted Fragment (FAU) Formulas in the almost uninterpreted fragment [14] are defined over a vocabulary that consists of the usual interpreted symbols of Linear Integer Arithmetic (LIA), equality and bit-vectors, extended with uninterpreted constant, function and relation symbols. In this work, we will not use bit-vectors. We recall that LIA includes constant symbols (e.g., $1, 2, \dots$), function symbols of linear arithmetic (e.g., “ $+$ ”, but not multiplication), and relation symbols (e.g., “ \leq ”), all of which are interpreted by the theory, which includes all formulas over this vocabulary that are satisfied by the integers. A formula (over the extended vocabulary) is in the *essentially uninterpreted* fragment if variables are restricted to appear as arguments to uninterpreted function or relation symbols. The *almost uninterpreted* fragment also allows variables to appear in inequalities in a restricted way (for example, inequalities between a variable and a ground term are allowed). For example $\forall x : \text{int}. x + y \leq z$, is not in the fragment, since the variable x appears under the interpreted relation \leq . However $\forall x : \text{int}. f(x) + y \leq z$ is allowed, as is $\forall x : \text{int}. x \leq y$.

The *finite* almost interpreted fragment (FAU) is defined as the set of almost interpreted formulas that are stratified as defined in [14]. Satisfiability of FAU modulo the theory is decidable. In particular, in [14] a set of groundings is defined that is sufficient for completeness. In FAU, this set is finite, which implies decidability. Moreover, it implies that every satisfiable formula has a model in which the universes of the uninterpreted sorts are finite. This is useful for providing counterexamples. The FAU fragment also subsumes the array property fragment described in [6].

Of particular importance for our purposes, the SMT solver Z3 [10] is a decision procedure for the FAU fragment. This is

because its model-based quantifier instantiation procedure guarantees to eventually generate every grounding in the required set. This gives us a rich language in which to express our verification conditions, without sacrificing decidability. The subject of this paper is how to structure the proofs of implementations of distributed protocols such that the verification conditions fall into this fragment.

4 Modular Proofs

In this section we describe an illustrative modular reasoning system, using a very simple procedural language as a model of MDL. This system is not as rich as the system actually used in the IVy tool but is sufficient to capture the proof strategies we apply here.

4.1 A model language

Let Σ be a vocabulary of non-logical symbols. The set of propositions (terms of Boolean sort) over Σ is denoted \mathcal{P} .

4.1.1 Statements

The statements in our model language are defined as follows:

Definition 4.1. Let \mathcal{N} be a set of *procedure names* and $\mathcal{V}_P \subset \Sigma$ a set of program variables. The *program statements* \mathcal{S} are defined by the following grammar:

$$\begin{aligned} \mathcal{S} ::= & c : \tau := t : \tau \mid (\mathcal{S}; \mathcal{S}) \mid \text{while } p \mathcal{S} \\ & \mid \text{if } p \mathcal{S} \mathcal{S} \mid \text{call } n \mid \text{skip} \end{aligned}$$

where c is in \mathcal{V}_P , t a term over Σ , τ a sort, $p \in \mathcal{P}$, and $n \in \mathcal{N}$.

The mutable program variables are a subset of the logical symbols. The statements have the expected semantics. For now, program variables c are restricted to logical constants, and can only be assigned terms t of first-order types. We will relax this restriction in Section 4.6.

4.1.2 Procedures and Modules

The *Hoare triples* \mathcal{H} are denoted $\{\varphi\} \sigma \{\psi\}$, where $\varphi, \psi \in \mathcal{P}$ and $\sigma \in \mathcal{S}_{\mathcal{N}}$. Our notion of procedure definition is captured by the following definition:

Definition 4.2. A *context* is a partial function from \mathcal{N} to \mathcal{H} . A context is denoted by a comma-separated list of *procedure definitions* of the form $n := H$, where $n \in \mathcal{N}$ and $H \in \mathcal{H}$, such that the names n are unique.

Intuitively, a context is a collection of procedure definitions with corresponding pre/post specifications. In MDL, the precondition φ of a procedure is introduced by the “requires” keyword and the postcondition by “ensures”.

A *module* is a procedural program that exports procedure definitions to its environment and has a determined set of initial states. In the sequel, if f is a partial function, we will write $\text{pre}(f)$ for its pre-image and $\text{img}(f)$ for its image. If P is a context, we will write $\text{called}(P)$ for the set of names n such that “call n ” occurs in P .

Definition 4.3. A *module* is a tuple (P, E, I, Q) , where:

- P is a context.
- $E \subseteq \text{pre}(P)$ is the set of *exports*.
- $I \subseteq \mathcal{P}$ is the *initial condition* of the module.
- $Q \subseteq \mathcal{P}$ is the *module invariant*.

That is, P gives a set of procedure definitions with pre/post specifications, E gives the subset of these definitions that is exported to the environment, I is a set of predicates that are true in the module’s initial state, and $\wedge Q$ is an inductive invariant that holds between calls to exported procedures. In the sequel, we often use I to denote $\wedge I$ and Q to denote $\wedge Q$. We will write P_M , E_M , I_M , Q_M , respectively, for the components of module M . In MDL, the invariant Q_M is given by the set of “invariant” declarations in the module.

We define the following operations on modules:

Definition 4.4. Let $M = (P, E, I, Q)$ and $M' = (P', E', I', Q')$ be modules such that $\text{pre}(P_M) \cap \text{pre}(P_{M'}) = \emptyset$ and $\Sigma(I_M) \cap \Sigma(I_{M'}) \cap \mathcal{V}_P = \emptyset$. The *composition* of M' and M , denoted $M' + M$, is $(P \cup P', E \cup E', I \cup I', Q \cup Q')$.

Definition 4.5. For a module $M = (P, E, I, Q)$ and a set $E' \subseteq \mathcal{N}$, the *restriction* of M to E' , denoted $M \downarrow E'$, is the module $(P, E \cap E', I, Q)$.

4.2 Axiomatic Semantics

We write $P \vdash_T \{\varphi\} \sigma \{\psi\}$ to denote the judgment that, *assuming* context P and background theory T , if σ starts at a T -model satisfying φ and terminates in a T -model, then this model satisfies ψ . In derivation rules, we will drop the theory T if it is the same for all judgments.

The axiomatic semantics of the statements of our language is given by the standard rules of Hoare logics:

$$\begin{array}{c} T \models (\varphi' \Rightarrow \varphi) \wedge (\psi \Rightarrow \psi') \\ P \vdash_T \{\varphi\} \sigma \{\psi\} \\ \hline \begin{array}{c} P \vdash_T \{\varphi'\} \sigma \{\psi'\} \\ P \vdash \{\varphi\} \sigma \{\psi\} \\ P \vdash \{\psi\} \sigma' \{\rho\} \\ \hline P \vdash \{\varphi\} (\sigma; \sigma') \{\rho\} \\ P \vdash \{\varphi \wedge p\} \sigma \{\varphi\} \\ \hline P \vdash \{\varphi\} \text{ while } p \sigma \{\varphi \wedge \neg p\} \\ P \vdash \{\varphi \wedge p\} \sigma \{\psi\} \\ P \vdash \{\varphi \wedge \neg p\} \sigma' \{\psi\} \\ \hline P \vdash \{\varphi\} \text{ if } p \sigma \sigma' \{\psi\} \\ \\ P \vdash \{\varphi[t / c]\} c := t \{\varphi\} \\ \\ P \vdash \{\varphi\} \text{ skip } \{\varphi\} \\ P \vdash \{\varphi\} \sigma \{\psi\} \\ \hline n := \{\varphi'\} \sigma \{\psi'\}, P \vdash \{\varphi \wedge \varphi'\} \text{ call } n \{\psi \wedge \psi'\} \end{array} \end{array}$$

Is the inlining rule the only one for calls?

Seems unpleasant to be so unmodular.

Oh, another rule comes later. Warn us here?

The first is the so-called “rule of consequence”. The remainder, respectively, give the semantics of sequential composition, while loops, conditionals, assignments and “skip”. The last rule is the Inline rule that gives the semantics of non-recursive procedure calls: any fact that can be proved about the body of procedure n in a given context can be used at a call site of n . Notice that we must still satisfy any specified pre-condition φ' and may use the specified post-condition ψ' . In effect, this allows us to inline a procedure definition at a call site. This is relatively complete for non-recursive programs, which include the examples we treat here.

We will write $I; P \vdash M$ to represent the judgment that, in context P , if formulas I hold initially, then module M maintains its invariant and satisfies its pre/post specifications. It is assumed that the environment only calls M 's exported procedures, and otherwise never modifies its program variables. We elide I or P if they are empty sets. The axiomatic semantics of modules is given by the following Module rule:

$$\frac{\begin{array}{c} I, I_M \models Q_M \\ \text{for } n := \{\varphi\} \sigma \{\psi\} \text{ in } P_M: \\ P, P_M \vdash \{\varphi\} \sigma \{\psi\} & \text{if } n \in \text{called}(P_M) \\ P, P_M \vdash \{Q_M \wedge \varphi\} \sigma \{Q_M \wedge \psi\} & \text{if } n \in E_M \end{array}}{I; P \vdash M}$$

In this rule, Q_M is an inductive invariant that holds between calls to exported procedures. It must hold in the initial states and be preserved by all exported procedures. In addition, internally called procedures must satisfy their specifications *without* assuming the invariant, since the invariant may be violated during execution of the module's procedures.

4.3 Rules for Decidable Decomposition

The rules defined in Section 4.2 provide the full axiomatic semantics. In particular, they enable to verify a program which consists of multiple modules, say M_1, \dots, M_n , and exports procedures E to its environment, by proving $\vdash (M_1 + \dots + M_n) \downarrow E$. However, the generated verification conditions may be undecidable. In this section, we provide inference rules that are derived from the semantics and enable to decompose the proof to obtain decidable verification conditions. For simplicity, we only present the rules needed for the Toy Leader Election example. Our implementation includes more flexible rules that are similar in spirit.

First, we introduce a rule that allows to verify a procedure call without inlining the procedure's body, by using the assumption that the procedure satisfies its pre/post specification at the call site:

Theorem 4.6. *The following rule Call can be derived:*

$$n := \{\varphi\} \sigma \{\psi\}, P \vdash \{\varphi\} \text{ call } n \{\psi\}$$

The next rules allow to verify the composition of modules by verifying the individual modules. We begin by defining some notation.

Definition 4.7. The *callset* $\text{calls}(M, P)$ of module M in context P is the least set of procedures n such that either n is exported from M , or n is in $\text{pre}(P_M \cup P)$ and n is called from $\text{calls}(M, P)$. Formally, $\text{calls}(M, P) = \text{LFP.} \lambda X. E_M \cup X \cup (\text{called}(X) \cap \text{pre}(P_M \cup P))$. The *refset* $\text{ref}(M, P)$ is the subset of \mathcal{V}_P occurring in $\text{calls}(M, P)$ or in I_M . The *modset* $\text{mod}(M, P)$ is the subset of \mathcal{V}_P assigned in $\text{calls}(M, P)$, or occurring in I_M . Module M is said to *interfere* with module M' in context P , denoted $M, P \rightsquigarrow M'$, if $\text{mod}(M, P) \cap \Sigma(Q_{M'}) \neq \emptyset$ or if $\text{called}(\text{calls}(M, P)) \cap \text{pre}(P_{M'}) \not\subseteq E_{M'}$.

In other words, M interferes with M' if it either modifies a variable occurring in the invariant of M' , or if it calls an internal procedure of M' . Module M can interfere with M' directly, or by calling procedures defined in the context P .

To export an invariant from one module to another, we introduce two notations. If M is a module and Γ is a set of formulas, we say $M[\Gamma]$ is the module that results from conjoining $\wedge \Gamma$ to the postcondition of every exported procedure of M . On the other hand, $[\Gamma]M$ results from conjoining $\wedge \Gamma$ to the precondition of every exported procedure of M .

Now we can derive a compositional rule that allows us to verify a service M' layered on a service M , while *assuming* the invariants and pre/post specifications of M .

Theorem 4.8. *The following rule Layer can be derived:*

$$\frac{\begin{array}{c} I; P \vdash M & M' \downarrow E, P \not\rightsquigarrow M \\ I, \Theta; P, P_{M[\Gamma]} \vdash [\Gamma]M' & M \downarrow E, P \not\rightsquigarrow M' \\ \hline I; P \vdash (M + M') \downarrow E & \Gamma \subseteq Q_M \\ & \Theta \subseteq I_M \end{array}}{}$$

Sloppy to call individual Ignore for the moment the expressions in square brackets. The rule states that, to verify M' layered on M , we first verify M , then verify M' assuming the proved specifications of the exported procedures of M . Intuitively, this works because external calls to one module cannot invalidate the invariant of the other. Note the asymmetry, however. Module M must be proved in context P , which means that it contains no calls to procedures outside of P , and in particular, no callbacks into M' . At a call-back, the invariant of M' would not hold, violating the precondition under which M' is verified. Technically, this rule can be derived by annotating every statement of M' with the invariants of M . The rule also allows us to use initial conditions of M .

The bracketed expressions allow us to assume the proved invariants of M when proving M' . We do this by assuming these invariants on entry to every exported procedure of M' and on exit of every exported procedure of M .

4.4 Ghost modules and slicing

Definition 4.9. If P is a context, the *slice* of P , denoted $\text{slice}(P)$ is the set of procedure definitions which contains $n := \{\text{true}\} \text{skip} \{\text{true}\}$ for each $n := \{\varphi\} \sigma \{\psi\}$ in P . If M is a module, $\text{slice}(M)$ denotes $(\text{slice}(P_M), E_M, \emptyset, \emptyset)$.

991 The following derived rule can be used to slice out a “ghost”
 992 module that is used only for the purpose of the proof. We
 993 say a module M is *invisible* to M' in context P , denoted
 994 $M, P \not\hookrightarrow M'$, if $M, P \not\rightsquigarrow M'$ and $\text{mod}(M, P) \cap \text{ref}(M', P) = \emptyset$
 995 and I_M is \mathcal{V}_P -conservative and every exported procedure of
 996 M' terminates in context P , starting in all states satisfying
 997 its precondition.
 998

999 **Theorem 4.10.** *The following inference rule can be derived:*

$$\frac{P \vdash (M + M') \downarrow E_{M'}}{P \vdash (\text{slice}(M) + M') \downarrow E_{M'}} \quad M, P \not\hookrightarrow M'$$

inconsistent capitalization

1000 To prove termination for the examples presented here,
 1001 it suffices to verify that there is no recursion and that P_M
 1002 contains no “while” statements IVy supports proof of termina-
 1003 tion using a ranking). We must also show that every model
 1004 of the theory has an extension to the program variables sat-
 1005 isfying I_M . In practice we must prove this using Theorem 3.2,
 1006 which means that I_M must be a conjunction of a sequence
 1007 of definitions.
 1008

1009 The invisible module M can be used like a lemma in the
 1010 proof of M' . That is, we make use of its properties and then
 1011 discard it, as we did with the abstract protocol model in Toy
 1012 Leader Election.
 1013

4.5 Segregating theories

1014 To allow us to segregate theories, we add one derived rule
 1015 Theory to our system:

1016 **Theorem 4.11.** *The following inference rule can be derived:*

$$\frac{T, T' \models T'' \quad P \vdash_{T \cup T''} \{\varphi\} \sigma \{\psi\}}{P \vdash_{T \cup T'} \{\varphi\} \sigma \{\psi\}}$$

1017 In other words, what can be proved in a weaker theory can
 1018 be proved in a stronger theory. This allows us, for example,
 1019 to replace the theory of arithmetic with the theory of linear
 1020 order, or to drop function definitions that are not needed
 1021 in a given module. In Toy Leader Election, for example, we
 1022 dropped the theory LIA and the definition of majority when
 1023 verifying the abstract model and the implementation, but
 1024 used them when verifying nset.
 1025

4.6 Language extensions

1026 In this section we introduce some useful extensions to the
 1027 basic language which, while straightforward, cannot be de-
 1028 tailed here due to space considerations.
 1029

1030 Though we have modeled procedure calls as having no
 1031 parameters, it is straightforward to extend the language to in-
 1032 clude call-by-value with return parameters. In the following
 1033 we assume such an extension.
 1034

1035 We allow assignments of the form $f := \lambda x. e$, where f is
 1036 a function, since the resulting verification conditions can still
 1037 be expressed in first order logic [36]. In the compiled code,
 1038

1039 the resulting value of f is a function closure. The assignment
 1040 $f(a) := e$ is shorthand for $f := \lambda x. \text{if } x = a \text{ then } e \text{ else } f(x)$.
 1041

1042 We provide built-in theories for integers and bit vectors
 1043 (both with the usual arithmetic operators). In IVy, these theo-
 1044 ries are provided natively by Z3. Finite immutable arrays are
 1045 provided as an abstract datatype, with functions provided
 1046 for length and select, and axiomatically specified procedures
 1047 for update and element append. For each finite sort τ (such
 1048 as node in Toy Leader Election) the language provides an
 1049 array constant $\tau.\text{all}$ that contains all elements of sort τ . We
 1050 used this feature to define the notion of a majority of nodes
 1051 in Toy Leader Election.
 1052

4.7 Modeling network communication

1053 For simplicity, we will introduce only a model of a broadcast
 1054 datagram service, as used in Toy Leader Election. Other
 1055 services can be modeled similarly. For each sort μ of messages
 1056 transmitted on te network, we introduce an abstract relation
 1057 $\text{sent}(m : \mu)$ to represent the fact that message m has been
 1058 broadcast in the past. We add to the language a primitive
 1059 “send $m : \mu$ ” whose semantics is defined by the following
 1060 rule:
 1061

$$\frac{P \vdash \{\varphi\} \text{ sent}(m : \mu) := \text{true } \{\psi\}}{P \vdash \{\varphi\} \text{ send } m : \mu \{\psi\}}$$

1062 That is, the effect of “send m ” is to add message m to the
 1063 set of broadcast messages of sort μ . A module using net-
 1064 work services for sort μ exports a procedure “ recv_μ ” that is
 1065 called by the network. This procedure takes two parameters:
 1066 $p : \text{pid}$ to represent the receiving process id and $m : \mu$ to
 1067 represent the received message. We use $\vdash^M P$, where \mathcal{M} is
 1068 a collection of message sorts, to represent the judgment that
 1069 P satisfies its specifications when composed with a network
 1070 that handles messages of sorts in \mathcal{M} . This judgment can be
 1071 derived by the following Network rule:
 1072

$$\frac{\begin{array}{c} \text{sent}(m : \mu) \models_T \varphi \\ \vdash^M_T M \\ \vdash^{\mathcal{M}, \mu}_T M \downarrow (E_M \setminus \text{recv}_\mu) \end{array}}{\text{recv}_\mu \in E_M \quad P_M(\text{recv}_\mu) = \{\varphi\} \sigma \{\psi\}}$$

1073 In other words, we can assume that the system calls
 1074 $\text{recv}(p, m)$ only with messages m that have already been
 1075 broadcast. This yields a very weak network semantics, al-
 1076 lowing messages to be dropped, reordered and duplicated.
 1077 In MDL, we used the keyword “handles” to indicate which
 1078 procedures are used to handle received messages of a given
 1079 sort. The keyword “system” indicates a top-level module, to
 1080 which the above rule should be applied.
 1081

4.8 Proof of Toy Leader Election

1082 To illustrate the inference rules, we explain how to prove Toy
 1083 Leader Election by chaining them. Let M^1, M^2, M^3 be, respec-
 1084 tively, the modules nset, toy_protocol and toy_system. First,
 1085 we prove $\vdash_T M^1$ where theory T consists of the integer arith-
 1086 metic and array theories, and the definitions of the majority
 1087

1088 1089 1090 1091 1092 1093 1094 1095 1096 1097 1098 1099 1100

and member relations. Next we prove $P_{M^1[\Gamma]} \vdash [\Gamma] M^2$ (using the Module and Call rules). Here, Γ is the majority intersection invariant of nset. Notice we do not use theory T , to preserve decidability. We then add theory T using the Theory rule, and combine with the above using the Layer rule, to obtain $\vdash_T M^1 + M^2$. Then we prove:

$$I_{M^2}; P_{(M^1+M^2)[\Gamma']} \vdash [\Gamma'] M^3$$

Here Γ' is the invariant of toy_protocol used by toy_system. By separating the proof of this invariant, we avoided a function cycle. In this proof, we inline the procedures of the abstract model M_2 and use the Send rule to capture the semantics of message sending. Again using the Theory and Layer rules, we obtain $\vdash_T (M_1 + M_2 + M_3) \downarrow E_{M^3}$. We use the Slice rule to remove the ghost module, obtaining $\vdash_T (M_1 + \text{slice}(M_2) + M_3) \downarrow E_{M^3}$. Finally, the network rule hides the message handlers, giving the conclusion:

$$\vdash_T^M (M^1 + \text{slice}(M^2) + M^3) \downarrow \{\text{request_vote}\}.$$

This leaves request_vote as the only exported procedure, which is called in response to a client request to initiate the protocol. The result is a verified equivalent to the code of Figure 1. Note we don't write these steps explicitly. Rather, the tool infers them based the "uses", "open", "ghost", and "system" directives in code.

4.9 Concurrency and parametricity

Thus far, we have considered a purely sequential program that presents exported procedures to be called by its environment and assumes that each call terminates before the next call begins. This semantics is implicit in Definition 4.3 and the Module rule. In reality, however, calls with different values of the process id parameter p will be executed concurrently. We need to be able to infer that every concurrent execution is sequentially consistent, that is, it is equivalent to some sequential execution when only the local histories of actions are observed. To do this, we use Lipton's theory of left-movers and right-movers [28], in much the same way as is done in the IronFleet project [15]. Since this argument does not relate directly to the use of decidable theories, we only sketch it here.

First, we need to show that any two statements executed by two different processes, excepting "send" statements, are independent. To do this, we require that every exported procedure have a first parameter $p : \text{pid}$ (representing the process id). We verify statically that every program variable reference (after slicing the ghost modules) is of the form $f(p, \dots)$ where p is the process id parameter. Another way to say this is that all statements except send statements are "both-movers" in Lipton's terminology. Moreover, by construction, every call from the environment consists of an optional message receive operation, followed by a combination of both-movers and sends. Since receive is a right-mover and send is a left-mover, it follows that every call can

be compressed to an atomic operation, thus the system is sequentially consistent by construction.

When we compile a module to executable code, we take the parameter p as a fixed value given at initialization of the process. We use this constant value to partially evaluate all program variable references, thus allowing the compiled code to store the only the state of one process.

4.10 Verification conditions

We use the inference rules above to generate verification conditions (VC's) in the usual manner, as in tools such as ESC Java [13] and Dafny [25]. These are validity checks that result from the side conditions of the Module, Network, and Theory rules, and the standard rule of consequence. The verifier checks that each VC is in one of our decidable fragments (taking into account any built-in theories used) and issues a warning if it is not. The warning may exhibit, for example, a bad cycle in the function graph. In case a VC is determined to be invalid by the Z3 prover, the counter-model produced by Z3 is used to create a program execution trace that demonstrates the failure of the VC.

5 Evaluation

Redundant content now that you're unblinded?

To evaluate our methodology, we applied it to develop verified implementations of Raft [34] and Multi-Paxos [23]. We used the IVy [36] system, which uses Z3 [10] to check verification conditions. Both Raft and Multi-Paxos implement a centralized shared log abstraction, i.e. a write-once map from indices to values. This abstraction can be used to implement a distributed fault-tolerant service, using the state machine replication approach [38]. We verify the safety property of the log abstraction, i.e. that no two replicas ever disagree on the contents of the shared log.

5.1 Raft and Multi-Paxos

Raft The Raft protocol operates in a sequence of *terms*. In each term a leader is elected in a way that is similar to the toy leader election protocol presented in Section 2, and the leader then replicates its log on the other nodes.

We implemented the basic Raft protocol without log truncation. Roughly speaking, the implementation works as follows. The implementation initially elects a leader node which then accepts client requests, appends them to its in-memory log (no data is ever written to disk), and replicates its log over all replicas. The implementation continues to work correctly as long as a majority of nodes stay online and can communicate with each other. Should the leader fail, the implementation elects a new leader which takes over processing client requests. The absence of log truncation means that the data-structure representing the log locally on each node grows without bound. Features such as state transfer to slow replicas, crash recovery, or batching are not part of our implementation.

Our verified Raft implementation follows the modularity principles presented in Section 2. It consists of an abstract protocol ghost module and its correctness proof, a module for the system implementation, and two datatype modules for node sets with majority testing logs represented as built-in arrays. The implementation module is fairly straightforward, and similar to an implementation in an imperative C-like language. It contains function symbols that map node to various sorts that capture the node state.

The ghost module contains the core of the proof, and is where most of the verification effort was spent. It contains auxiliary relations, used both to avoid quantifier alternation cycles and to record information used in the inductive invariant. For example, while the concrete implementation contains a function $t : \text{node} \rightarrow \text{term}$, the ghost module abstracts it as a relation $t_g : \text{node}, \text{term}$, intended to capture $t_g(n, x) \equiv t(n) = x$. This avoids quantifier alternation from node to term.

Verification of the abstract protocol ghost module is done in EPR, with log indices and terms abstracted as linear orders. This module provides the key correctness invariant, i.e., consistency and persistence of committed entries. The verification of the implementation module is also done in EPR. However, the abstract protocol and the implementation use different quantifier alternation stratification orders, so a non-modular proof would lead to undecidability. The modules for nodesets and logs are verified in the FAU fragment.

Multi-Paxos Our approach to implementing Multi-Paxos and the features implemented are largely similar to that of Raft. We implemented the basic protocol. Log truncation, state transfer, batching, and recovery are not part of the implementation.

5.2 Verification Effort

The Raft verification took 2 person-months, including 1.5 person-months for verifying the abstract protocol, and 2 weeks for verifying the implementation. The code contains 560 SLOC of implementation and 200 SLOC of invariants and ghost code, for a proof-to-code ratio of 0.36. Obtaining the Multi-Paxos implementation from the abstract protocol (which was already proved in previous work [35]) took approximately one week of work, reusing some of the work done for Raft. The code consists of 330 SLOC of implementation and 266 SLOC of invariants and ghost code. In total, the proof to code ratio is of about 0.8 (counting invariants and ghost code as proof). IVy successfully discharges all VCs of both Multi-Paxos and Raft in a few minutes on a conventional laptop. During the development, IVy quickly produced counterexamples to induction and displayed them graphically, which greatly assisted in the verification process.

For comparison, IronFleet's IronRSL [15] implementation is part of a verification effort of 3.7 person-year (including another verified protocol), and VC checking was outsourced to

Why two?

System	Throughput	Get	Put
IVy-Raft	550.1	84.4	85.1
IVy-Multi-Paxos	554.2	83.9	84.0
vard	624.2	79.8	80.0
etcd	2943.5	16.5	17.1

Table 1. Throughput (requests/s) and average request latency (ms) of gets and puts to the key-value stores.

a cloud to obtain verification times acceptable for interactive use.

As evidenced by its larger code-base, the IronRSL implementation has more features than the IVy implementations presented: log truncation, batching, and state-transfer are part of IronRSL but not of the IVy implementations (however both IronRSL and the IVy implementations do not support crash recovery). Moreover, more properties are verified compared to the IVy implementations: IronFleet verifies the network serialization and deserialization code, some liveness properties, and the model includes resource bounds (e.g. integer overflows), while those are not verified in the IVy implementations.

Excluding generic IronFleet libraries such as verified serialization code and liveness proofs, which are not verified in the IVy implementations, IronRSL consists of roughly 3000 implementation SLOC and 12000 proof SLOC. This gives a proof/code ratio of about 4.

The Verdi proof of the Raft protocol [41] consists of 50000 lines of proof for 530 lines of implementation and required several person-years. Among those 50000 lines, at least 95% are specific to the Raft protocol, while less than 5% can be counted as generic and reusable for other protocols. The proof is done in the Coq [4] proof assistant and has a smaller trusted code base than IVy.

5.3 Verified System's Performance

To evaluate the performance of our verified systems, we developed a key-value store that replicates data across multiple nodes using either Raft or Multi-Paxos. The implementations extracted from IVy handle all internal communication, while the (unverified) key-value component handles communication with clients and data storage; it consists of 785 SLOC of C++. We benchmarked the performance of these systems against that of vard [40], a verified key-value store, as well as etcd [9], an unverified, production-quality key-value store. Both vard and etcd use Raft for consensus.

We benchmarked all systems on a cluster of three Amazon EC2 t2.small nodes, each with 1 CPU core and 2GB of RAM. We then started 50 closed-loop client threads on a fourth t2.small node. The threads sent 10000 requests (a randomized 50/50 mix of GETs and PUTs). vard and etcd both persist data to disk, while our key-value store does not; therefore we modified the vard shim to stop it from writing to disk, and we configured etcd to use a RAM-disk.

1321 Our results are summarized in Table 1. Both our verified
 1322 key-value stores achieve similar performance to `vard`, while
 1323 all three verified stores are roughly 5× slower than `etcd`. We
 1324 have not significantly optimized our implementations, so it
 1325 is unsurprising that a production-quality system achieves
 1326 better performance.

1327

6 Related Work

1329

Fully Automatic Verification Fully automatic push-button verification is usually beyond reach because of undecidability. Bounded checking is successfully used in systems like Alloy [18] and TLA+ [24] to check correctness of designs up to certain number¹⁶ of nodes. This is useful, due to the observation that most bugs occur with small number¹⁷ of nodes. However as observed by Amazon [32] and others it is hard to scale these methods even for very few, e.g., 3 nodes. Also many of the interesting bugs occur in the implementation. One interesting approach to obtain decidability of fully automatic verification is via limiting the class of programs [5, 20–22]. Another approach is to use sound and incomplete procedures for deduction and invariant search for logics that combine quantifiers and set cardinalities [39]. However, distributed systems of the level of complexity we consider here (i.e., Raft, Multi-Paxos) are beyond¹⁸ reach of this technique. Another direction, explored in [2], is to verify limited properties (e.g., for absence of deadlocks), using a sound but incomplete decidable check. None of the state-of-the-art techniques for fully automatic verification can prove properties such as consistency for systems implementations. Moreover, when automatic methods fail, the user is struck without any solution.

1353

1354

Interactive Verification Recent projects such as Verdi [40] and IronFleet [15] demonstrate that distributed systems can be proved all the way from the design to an implementation. The DistAlgo project [7, 29] develops programming methodologies for interactively verifying distributed systems, based on the TLA+ proof system [8]. However, interactive verification requires tremendous human efforts and skills and thus has limited success so far. Our work can be considered as an attempt to understand how much automation is achievable using modularity. We argue that invariants provide a reasonable way to interact with verification systems since one does not need to understand how decision procedures such as SMT work¹⁹. In this work we express invariants in first-order logic. In the future, it may be possible to develop high-level specification approaches.

1368

1369

1370

1371

1372

1373

1374

1375

1376

1377

1378

1379

1380

1381

1382

1383

1384

1385

1386

1387

1388

1389

1390

1391

1392

1393

1394

1395

1396

1397

1398

1399

1400

1401

1402

1403

1404

1405

1406

1407

1408

1409

1410

1411

1412

1413

1414

1415

1416

1417

1418

1419

1420

1421

1422

1423

1424

1425

1426

1427

1428

1429

1430

which develops a decidable fragment that allows some arithmetic with function symbols and cardinality constraints. Compared to these works, our approach considers a more general setting of asynchronous communication, and uses more restricted and mainstream decidable fragments that are supported by existing theorem provers. This is enabled by our use of modularity. Our approach of applying modularity to obtain decidability will benefit and become more powerful as more expressive decidable logics are developed and supported by efficient solvers.

Recently, [35] showed how to verify abstract protocols of the Paxos family in EPR, and [31] used it for cache coherence. In this work we use EPR to reason about abstract protocols but also go beyond the protocol level to verify an actual efficient system implementation. The gap between the abstract protocol and the system implementation presents a major technical challenge. We bridge this gap by the novel use of modularity to obtain decidability, as well as the use of the FAU fragment [14], which is more expressive than EPR.

Modularity in Verification The utility of modularity for simplified reasoning was already recognized in the seminal works of Hoare and Dijkstra (e.g., [17]). Proof assistants such as Isabelle/HOL [33] and Coq [4] provide various modularity mechanisms. Indeed, existing deductive verification engines such as Dafny [25] employ modularity to simplify reasoning in a way that is similar to ours. In this landscape, our chief novelties are the use of modularity for decidability, and a methodology for modular, decidable reasoning about distributed systems. We note that unlike Dafny, Ivy performs syntactic checks to ensure that verification conditions are in decidable logics. Thus, the user either receives an error message and corrects the specification or can be assured that the verification problem is solvable. Our evaluation demonstrates that our methodology is useful for verifying complex distributed systems, and that it drastically reduces the verification effort.

7 Conclusion

Modularity is well recognized as a key to scalability of systems. This paper shows that modularity enables decidability of reasoning about real implementations of distributed protocols. At the implementation level, distributed protocols involve arithmetic, unbounded sets of processes and unbounded data structures. For this reason, we might expect that reasoning about these systems would require the use of undecidable logics. We have seen however, that by a fairly simple modular decomposition, we can separate the proof into lemmas that reside in decidable fragments, which in turn can make the use of automated provers more predictable and transparent.

References

- [1] Francesco Alberti, Silvio Ghilardi, and Elena Pagani. 2016. Counting Constraints in Flat Array Fragments. In *Automated Reasoning*. Springer,

- 1431 Cham, 65–81.
 1432 [2] Alexander Bakst, Klaus von Gleissenthall, Rami Gökhan Kici, and
 1433 Ranjit Jhala. 2017. Verifying distributed programs via canonical se-
 1434 quentialization. *PACMPL 1, OOPSLA* (2017), 110:1–110:27. <https://doi.org/10.1145/3133934>
 1435 [3] Josh Berdine, Cristiano Calcagno, and Peter W. O’Hearn. 2004. A
 1436 Decidable Fragment of Separation Logic. In *FSTTCS 2004: Foundations
 1437 of Software Technology and Theoretical Computer Science, 24th International
 1438 Conference, Chennai, India, December 16–18, 2004, Proceedings*. 97–109.
 1439 [4] Yves Bertot and Pierre Castérán. 2004. *Interactive Theorem Proving and
 1440 Program Development - Coq’Art: The Calculus of Inductive Constructions*.
 1441 Springer. <https://doi.org/10.1007/978-3-662-07964-5>
 1442 [5] Roderick Bloem, Swen Jacobs, Ayrat Khalimov, Igor Konnov, Sasha
 1443 Rubin, Helmut Veith, and Josef Widder. 2015. *Decidability of Pa-
 1444 rameterized Verification*. Morgan & Claypool Publishers. <https://doi.org/10.2200/S00658ED1V01Y201508DCT013>
 1445 [6] Aaron R. Bradley, Zohar Manna, and Henny B. Sipma. 2006. What’s
 1446 Decidable About Arrays?. In *Verification, Model Checking, and Abstract
 1447 Interpretation, 7th International Conference, VMCAI 2006, Charleston,
 1448 SC, USA, January 8–10, 2006, Proceedings*. 427–442. https://doi.org/10.1007/11609773_28
 1449 [7] Saksham Chand, Yanhong A. Liu, and Scott D. Stoller. 2016. Formal Ver-
 1450 ification of Multi-Paxos for Distributed Consensus. In *FM 2016: Formal
 1451 Methods: 21st International Symposium, Limassol, Cyprus, November
 1452 9–11, 2016, Proceedings 21*. Springer, 119–136.
 1453 [8] Kaustuv Chaudhuri, Damien Doligez, Leslie Lamport, and Stephan
 1454 Merz. 2010. The TLA+Proof System: Building a Heterogeneous Veri-
 1455 fication Platform. In *Proceedings of the 7th International Colloquium
 1456 Conference on Theoretical Aspects of Computing (ICTAC’10)*. Springer-
 1457 Verlag, 44–44.
 1458 [9] CoreOS 2014. etcd: A highly-available key value store for shared
 1459 configuration and service discovery. (2014). <https://github.com/coreos/etcd>.
 1460 [10] Leonardo de Moura and Nikolaj Bjørner. 2008. Z3: An Efficient SMT
 1461 Solver. In *Tools and Algorithms for the Construction and Analysis of
 1462 Systems, 14th International Conference, TACAS 2008, Held as Part of the
 1463 Joint European Conferences on Theory and Practice of Software, ETAPS
 1464 2008, Budapest, Hungary, March 29–April 6, 2008. Proceedings (Lecture
 1465 Notes in Computer Science)*, Vol. 4963. Springer, 337–340.
 1466 [11] Cezara Dragoi, Thomas A. Henzinger, Helmut Veith, Josef Widder,
 1467 and Damien Zufferey. 2014. A Logic-Based Framework for Verifying
 1468 Consensus Algorithms. In *International Conference on Verification,
 1469 Model Checking, and Abstract Interpretation*. Springer, 161–181.
 1470 [12] Cezara Dragoi, Thomas A. Henzinger, and Damien Zufferey. 2016.
 1471 PSync: A Partially Synchronous Language for Fault-Tolerant Dis-
 1472 tributed Algorithms. *ACM SIGPLAN Notices* 51, 1 (2016), 400–415.
 1473 [13] Cormac Flanagan, K. Rustan M. Leino, Mark Lillibridge, Greg Nelson,
 1474 James B. Saxe, and Raymie Stata. 2002. Extended Static Checking for
 1475 Java. In *Proceedings of the ACM SIGPLAN 2002 Conference on Program-
 1476 ming Language Design and Implementation (PLDI ’02)*. ACM, 234–245.
 1477 <https://doi.org/10.1145/512529.512558>
 1478 [14] Yeting Ge and Leonardo De Moura. 2009. Complete instantiation for
 1479 quantified formulas in satisfiability modulo theories. In *International
 1480 Conference on Computer Aided Verification*. Springer, 306–320.
 1481 [15] Chris Hawblitzel, Jon Howell, Manos Kapritsolis, Jacob R. Lorch, Bryan
 1482 Parno, Michael L. Roberts, Srinath T. V. Setty, and Brian Zill. 2015.
 1483 IronFleet: proving practical distributed systems correct. In *Proceedings
 1484 of the 25th Symposium on Operating Systems Principles, SOSP*. 1–17.
 1485 [16] Jesper G. Henriksen, Jakob L. Jensen, Michael E. Jørgensen, Nils Klær-
 1486 lund, Robert Paige, Theis Rauhe, and Anders Sandholm. 1995. Mona:
 1487 Monadic Second-Order Logic in Practice. In *Tools and Algorithms for
 1488 Construction and Analysis of Systems, First International Workshop,
 1489 TACAS*. 89–110.
 1490 [17] C. A. R. Hoare. 1972. Proof of correctness of data representations. 1, 4
 1491 (1972), 271–281.
 1492 [18] Daniel Jackson. 2006. *Software Abstractions: Logic, Language, and
 1493 Analysis*. The MIT Press.
 1494 [19] Gerwin Klein, June Andronick, Kevin Elphinstone, Gernot Heiser,
 1495 David Cock, Philip Derrin, Dhammika Elkaduve, Kai Engelhardt, Rafal
 1496 Kolanski, Michael Norrish, Thomas Sewell, Harvey Tuch, and Simon
 1497 Winwood. 2010. seL4: formal verification of an operating-system
 1498 kernel. *Commun. ACM* 53, 6 (2010), 107–115.
 1499 [20] Igor Konnov, Marijana Lazic, Helmut Veith, and Josef Widder. 2017. A
 1500 Short Counterexample Property for Safety and Liveness Verification
 1501 of Fault-Tolerant Distributed Algorithms. In *Proceedings of the 44th
 1502 ACM SIGPLAN Symposium on Principles of Programming Languages
 1503 (POPL 2017)*. ACM, 719–734.
 1504 [21] Igor Konnov, Helmut Veith, and Josef Widder. 2015. SMT and POR Beat
 1505 Counter Abstraction: Parameterized Model Checking of Threshold-
 1506 Based Distributed Algorithms. In *Computer Aided Verification*. Springer,
 1507 Cham, 85–102.
 1508 [22] Igor V. Konnov, Helmut Veith, and Josef Widder. 2015. What You
 1509 Always Wanted to Know About Model Checking of Fault-Tolerant
 1510 Distributed Algorithms. In *Perspectives of System Informatics - 10th
 1511 International Andrei Ershov Informatics Conference, PSI 2015, in Memory
 1512 of Helmut Veith, Kazan and Innopolis, Russia, August 24–27, 2015, Revised
 1513 Selected Papers (Lecture Notes in Computer Science)*, Manuel Mazzara
 1514 and Andrei Voronkov (Eds.), Vol. 9609. Springer, 6–21. https://doi.org/10.1007/978-3-319-41579-6_2
 1515 [23] Leslie Lamport. 1998. The Part-Time Parliament. *ACM Trans. Comput.
 1516 Syst.* 16, 2 (1998), 133–169. <https://doi.org/10.1145/279227.279229>
 1517 [24] Leslie Lamport. 2002. *Specifying Systems: The TLA+ Language and
 1518 Tools for Hardware and Software Engineers*. Addison-Wesley Longman
 1519 Publishing Co., Inc., Boston, MA, USA.
 1520 [25] K Rustan M Leino. 2010. Dafny: An automatic program verifier for
 1521 functional correctness. In *Logic for Programming, Artificial Intelligence,
 1522 and Reasoning*. Springer, 348–370.
 1523 [26] Xavier Leroy. 2009. Formal verification of a realistic compiler. *Commun.
 1524 ACM* 52, 7 (2009), 107–115.
 1525 [27] Harry R. Lewis. 1980. Complexity results for classes of quantificational
 1526 formulas. *J. Comput. System Sci.* 21, 3 (1980), 317 – 353.
 1527 [28] R.J. Lipton. 1975. Reduction: A method of proving properties of parallel
 1528 programs. *Commun. ACM* 18, 12 (1975), 717–721.
 1529 [29] Yanhong A. Liu, Scott D. Stoller, and Bo Lin. 2017. From Clarity to Effi-
 1530 ciency for Distributed Algorithms. *ACM Transactions on Programming
 1531 Languages and Systems* 39, 3 (July 2017).
 1532 [30] P. Madhusudan, Gennaro Parlato, and Xiaokang Qiu. 2011. Decid-
 1533 able logics combining heap structures and data. In *Proceedings of the
 1534 38th ACM SIGPLAN-SIGACT Symposium on Principles of Programming
 1535 Languages, POPL 2011, Austin, TX, USA, January 26–28, 2011*. 611–622.
 1536 [31] Kenneth L. McMillan. 2016. Modular specification and verification
 1537 of a cache-coherent interface. In *2016 Formal Methods in Computer-
 1538 Aided Design, FMCAD 2016, Mountain View, CA, USA, October 3–6,
 1539 2016*, Ruzica Piskac and Muralidhar Talupur (Eds.). IEEE, 109–116.
 1540 <https://doi.org/10.1109/FMCAD.2016.7886668>
 1541 [32] Chris Newcombe, Tim Rath, Fan Zhang, Bogdan Munteanu, Marc
 1542 Brooker, and Michael Deardeuff. 2015. How Amazon web services
 1543 uses formal methods. *Commun. ACM* 58, 4 (2015), 66–73.
 1544 [33] Tobias Nipkow, Lawrence C. Paulson, and Markus Wenzel. 2002. *Is-
 1545 abelle/HOL: A Proof Assistant for Higher-Order Logic*. Vol. 2283. Springer
 1546 Science & Business Media.
 1547 [34] Diego Ongaro and John K. Ousterhout. 2014. In Search of an Under-
 1548 standable Consensus Algorithm. In *2014 USENIX Annual Technical Con-
 1549 ference, USENIX ATC ’14, Philadelphia, PA, USA, June 19–20, 2014*. 305–
 1550 319. <https://www.usenix.org/conference/atc14/technical-sessions/presentation/ongaro>

1486
 1487
 1488
 1489
 1490
 1491
 1492
 1493
 1494
 1495
 1496
 1497
 1498
 1499
 1500
 1501
 1502
 1503
 1504
 1505
 1506
 1507
 1508
 1509
 1510
 1511
 1512
 1513
 1514
 1515
 1516
 1517
 1518
 1519
 1520
 1521
 1522
 1523
 1524
 1525
 1526
 1527
 1528
 1529
 1530
 1531
 1532
 1533
 1534
 1535
 1536
 1537
 1538
 1539
 1540

- 1541 [35] Oded Padon, Giuliano Losa, Mooly Sagiv, and Sharon Shoham. 2017.
 1542 Paxos Made EPR: Decidable Reasoning About Distributed Protocols.
 1543 *Proc. ACM Program. Lang.* 1, OOPSLA, Article 108 (Oct. 2017), 31 pages.
<https://doi.org/10.1145/3140568>
- 1544 [36] Oded Padon, Kenneth L. McMillan, Aurojit Panda, Mooly Sagiv, and
 1545 Sharon Shoham. 2016. Ivy: safety verification by interactive generalization. In *Proceedings of the 37th ACM SIGPLAN Conference on Programming Language Design and Implementation, PLDI 2016, Santa Barbara, CA, USA, June 13–17, 2016*. 614–630.
- 1546 [37] F. Ramsey. 1930. On a problem in formal logic. In *Proc. London Math. Soc.*
- 1547 [38] Fred B. Schneider. 1990. Implementing Fault-Tolerant Services Using
 1548 the State Machine Approach: A Tutorial. *ACM Computing Surveys (CSUR)* 22, 4 (1990), 299–319.
- 1549 [39] Klaus v. Gleissenthall, Nikolaj Bjørner, and Andrey Rybalchenko. 2016.
 1550 Cardinalities and Universal Quantifiers for Verifying Parameterized
 1551 Systems. In *Proceedings of the 37th ACM SIGPLAN Conference on Programming Language Design and Implementation (PLDI '16)*. ACM, 599–
 1552 613.
- 1553 [40] James R. Wilcox, Doug Woos, Pavel Panchevka, Zachary Tatlock, Xi
 1554 Wang, Michael D. Ernst, and Thomas E. Anderson. 2015. Verdi: a frame-
 1555 work for implementing and formally verifying distributed systems.
 1556 In *Proceedings of the 36th ACM SIGPLAN Conference on Programming Language Design and Implementation, Portland, OR, USA, June 15–17, 2015*. 357–368.
- 1557 [41] Doug Woos, James R. Wilcox, Steve Anton, Zachary Tatlock, Michael D.
 1558 Ernst, and Thomas E. Anderson. 2016. Planning for change in a formal
 1559 verification of the raft consensus protocol. In *Proceedings of the 5th ACM SIGPLAN Conference on Certified Programs and Proofs, Saint Petersburg, FL, USA, January 20–22, 2016*, Jeremy Avigad and Adam Chlipala
 1560 (Eds.). ACM, 154–165. <https://doi.org/10.1145/2854065.2854081>
- 1561
- 1562
- 1563
- 1564
- 1565
- 1566
- 1567
- 1568
- 1569
- 1570
- 1571
- 1572
- 1573
- 1574
- 1575
- 1576
- 1577
- 1578
- 1579
- 1580
- 1581
- 1582
- 1583
- 1584
- 1585
- 1586
- 1587
- 1588
- 1589
- 1590
- 1591
- 1592
- 1593
- 1594
- 1595

A Hoare logic derivation rules

These are the standard rules of Hoare logic that we apply:

$T \models (\varphi' \Rightarrow \varphi) \wedge (\psi \Rightarrow \psi')$	1596
$P \vdash_T \{\varphi\} \sigma \{\psi\}$	1599
<hr/>	1600
$P \vdash_T \{\varphi'\} \sigma \{\psi'\}$	1601
$P \vdash \{\varphi\} \sigma \{\psi\}$	1602
$P \vdash \{\psi\} \sigma' \{\rho\}$	1603
<hr/>	1604
$P \vdash \{\varphi\} (\sigma; \sigma') \{\rho\}$	1605
$P \vdash \{\varphi \wedge p\} \sigma \{\varphi\}$	1606
<hr/>	1607
$P \vdash \{\varphi\} \text{ while } p \sigma \{\varphi \wedge \neg p\}$	1608
$P \vdash \{\varphi \wedge p\} \sigma \{\psi\}$	1609
$P \vdash \{\varphi \wedge \neg p\} \sigma' \{\psi\}$	1610
$P \vdash \{\varphi\} \text{ if } p \sigma \sigma' \{\psi\}$	1611
<hr/>	1612
$P \vdash \{\varphi \langle t/c \rangle\} c := t \{\varphi\}$	1613
<hr/>	1614
$P \vdash \{\varphi\} \text{ skip } \{\varphi\}$	1615

The first is the so-called “rule of consequence”. The remainder, respectively, give the semantics of sequential composition, while loops, conditionals, assignments and “skip”.